

Rotation-Equivariant Variational Autoencoders for Representation Learning and Downstream Classification in Histopathology

Anonymous Authors

Abstract—At high magnification, cellular patterns in histopathology have no fixed orientation in the image, making rotation a relevant local symmetry for representation learning. This work compares conventional and continuous-SO(2)-equivariant denoising variational autoencoders (VAEs) trained on the same unlabeled patches. With both models frozen, reconstruction was evaluated on test patches, followed by two downstream tasks: three-class tissue recognition at patch level and five-class whole-slide image (WSI) diagnosis through multiple-instance learning. Mean reconstruction MAE was 0.0629 versus 0.0642 (conventional versus equivariant), but the paired confidence interval included zero. Tissue macro-F1 favored the equivariant representation at four of five label budgets; only the comparison at 500 labels per class had an adjusted interval excluding zero. On 23 test WSIs, macro-F1 was 0.389 versus 0.483, but the paired interval included zero.

Index Terms—histopathology, variational autoencoder, rotation equivariance, SO(2), whole-slide image classification, representation learning

I. INTRODUCTION

Digital histopathology converts glass slides into whole-slide images whose dimensions commonly reach gigapixel scale. These images contain many local tissue regions, but their diagnostic labels and detailed annotations require specialist review. Representation learning from unlabeled patches offers a way to use this image volume before fitting a smaller supervised model [1]–[3].

While orientation can be relevant to tissue architecture at lower magnification, cellular patterns at high magnification have no fixed orientation in the image. At this scale, a rotation changes the viewing orientation of a patch while preserving its tissue identity and diagnostic information. This rotational symmetry motivates incorporating rotation into the model’s architectural constraints. Ordinary convolutional neural networks (CNNs) use convolutions that are equivariant to translations: translating an input translates its feature map. Conventional CNNs usually learn rotational responses from rotated examples or separate filters. In the SO(2) VAE, steerable convolutions incorporate rotation symmetry through constrained filters, introducing a rotational inductive bias into the unsupervised representation learning performed by this VAE [4]–[6].

Rotation-equivariant classifiers have previously been studied for histopathology patch prediction [7]. VAEs incorporating symmetry have also modeled object pose and learned orientation-invariant cell representations [8], [9]. The comparison here uses patches from the UBC Ovarian Cancer Subtype

Classification and Outlier Detection (UBC-OCEAN) training resource [10]. A conventional residual denoising VAE and a VAE with steerable kernels constrained for the continuous planar rotation group SO(2) are trained without supervision [11], [12]. Both yield a spatial Gaussian latent map.

In this work, we aim to address two questions:

- 1) Does the implemented rotation-equivariant VAE yield lower reconstruction error on unseen histopathology patches than the conventional VAE?
- 2) Does the rotation-equivariant VAE representation yield higher classification macro-F1 under a shared supervised protocol: (a) for patch-level three-class tissue recognition from individual latent maps across label budgets, and (b) for WSI-level five-class diagnosis from bags of latent maps?

In this work, established steerable-kernel and field-aware operations are implemented in a continuous-SO(2) denoising VAE trained on unlabeled histopathology patches. Its encoder–decoder is compared with a conventional VAE in reconstruction; their frozen spatial latent maps are then compared in patch-level tissue recognition and WSI diagnosis with a patch classifier and an MIL classifier, respectively. The two models share image populations, latent dimensions, training budget and supervised assignments; uncertainty is assessed at WSI level. Reconstruction point estimates favored the conventional VAE, whereas tissue macro-F1 favored the equivariant representation at four of five label budgets. WSI diagnosis did not show a statistically significant difference between the representations.

II. RELATED WORK

A. Unsupervised Learning for Histopathology

Unsupervised learning uses unlabeled histopathology patches to learn representations that can subsequently support supervised prediction. Autoencoders learn these representations through reconstruction, while downstream tasks assess whether they retain information relevant to tissue recognition and slide-level diagnosis. Tellez *et al.* used neural image compression to learn patch representations without supervision before predicting gigapixel image labels. Their comparison included reconstruction, contrastive training and adversarial feature learning [3].

VAEs that incorporate image symmetries provide related precedents beyond histopathology. Nasiri and Bepler’s




Fig. 1. Rotation equivariance as a transformation law. The two paths from x to the lower-right field agree for an equivariant mapping f . The input rotation T_θ acts on the image; the corresponding output action T'_θ also rotates vector components when the output contains vector fields. Shapes and arrows are schematic rather than measured VAE features.

TARGET-VAE uses a translation- and rotation-equivariant encoder and a spatially equivariant generator to separate object pose from a semantic representation [8]. Burgess *et al.* introduced O2-VAE for cell and organelle shape profiling. It constructs orientation-invariant latent representations using equivariant encoder operations and pooling, and aligns the reconstruction with the input [9].

These studies motivate learning features without supervision and assessing their relevance through downstream tasks. Here, the VAEs are trained on unlabeled histopathology patches, evaluated in reconstruction, and frozen before their representations are used for patch-level tissue recognition and WSI-level diagnosis.

B. Rotation Equivariance

Ordinary CNNs build on translation-equivariant convolutions: shifting the input shifts the resulting feature map. Group-equivariant convolutional networks generalize this construction to other transformations, including rotations, by requiring layers to satisfy a prescribed transformation law for a selected group [4]. Figure 1 illustrates this relation: applying a rotation before an equivariant mapping corresponds to applying the prescribed rotation to its output. For vector-valued fields, this output transformation rotates both spatial locations and vector components. Steerable convolutions extend the construction to continuous rotations by restricting the kernel space according to the representations carried by the input and output fields [5], [6]. When the prediction task shares this symmetry, the restriction can reduce the set of functions the network can represent [13]. The architecture also changes parameter count and computational work; both are reported alongside model performance. Previous histopathology work evaluated rotation equivariance in supervised patch classifiers [7], [14].

C. Multi-Instance Learning

Whole-slide images can contain billions of pixels, making processing at native resolution in a single forward pass computationally impractical due to the high memory and computational requirements. A common approach therefore divides each WSI into patches, learns a local representation, and aggregates patch features for slide-level prediction. Attention-based multiple instance learning (MIL) provides a related supervised formulation in which a WSI is treated as a bag of instances and only the bag label is required [15]. CLAM combines attention pooling with instance-level clustering [16]. DSMIL combines contrastive patch representation learning with dual-stream aggregation [17], whereas TransMIL models relations among patch instances during aggregation [18]. Here, MIL tests the frozen patch representations under the same WSI classification protocol. Figure 2 illustrates the slide-to-bag formulation.

III. MATERIALS AND METHODS

A. UBC-OCEAN dataset

The UBC-OCEAN ovarian-cancer training resource contains 538 slide images: 513 conventional whole-slide images (WSIs) and 25 tissue microarray (TMA) images [10]. A WSI shows a large tissue section; a TMA shows multiple small tissue cores in one image. The resource describes WSIs as large, approximately $20\times$ scans and TMAs as smaller, approximately 4000×4000 -pixel, $40\times$ images. This scale and layout difference motivated using only the 513 non-TMA WSIs in the present experiment. Figure 3 shows a representative image of each type and their original pixel dimensions. For an image of width W and height H , a non-overlapping 256-pixel grid contains approximately $\lfloor W/256 \rfloor \lfloor H/256 \rfloor$ full patches before foreground filtering. At the dimensions quoted by the resource, a 4000×4000 image yields about 225 full patches, whereas a 100000×50000 image yields about 76,000; the latter is an illustrative large-WSI scale, not a count for an individual study slide. The resource also provides 152 supplemental color masks paired with a subset of those WSIs. These masks mark selected tumor, stroma and necrosis regions; they do not annotate every tissue pixel. The other 361 WSIs have no supplemental mask. The diagnostic label of each image is one of clear cell carcinoma (CC), endometrioid carcinoma (EC), high-grade serous carcinoma (HGSC), low-grade serous carcinoma (LGSC), or mucinous carcinoma (MC).

Each WSI was divided into non-overlapping RGB 256×256 -pixel patches. A full WSI has one diagnosis label, so thousands of patches extracted from that slide share the same slide-level label. Local tissue labels, in contrast, come only from eligible regions of the supplemental masks. Splits were made by WSI: no patches from one slide crossed train, validation or test boundaries.

Figure 3 shows two full-resolution patches beside their source WSI thumbnail, a partial supplemental mask on a second WSI, and a TMA. The thumbnail makes the slide-to-patch scale change visible; model inference used crops




Fig. 2. Multiple-instance learning for WSI diagnosis. Many patches extracted from one slide form a bag. Their encoded features are aggregated to predict one whole-slide diagnosis. The six thumbnail regions illustrate the slide-to-bag construction; the ellipsis represents the many remaining instances. In this study, patch features are obtained from frozen VAE posterior means, and supervision is provided by the WSI diagnosis label.

from the full-resolution WSI. The mask marks selected tissue regions rather than the full image.

B. Patch Selection and Tissue Annotations

Foreground selection for reconstruction and WSI diagnosis was derived from the image rather than from human masks. Each WSI thumbnail was converted to HSV, and Otsu’s threshold was applied to the saturation channel [20]. A full-resolution grid cell was retained when more than 60% of its projected thumbnail area belonged to foreground. The resulting coordinates were used for VAE development, reconstruction and WSI diagnosis without consulting the supplemental tissue masks.

Supplemental masks were used only to construct the tissue-recognition population. Red, green and blue mask regions were interpreted as tumor, stroma and necrosis. The masks are not exhaustive; an unpainted region cannot be treated as a negative class. A patch was therefore eligible for tissue classification only when at least 10% of its area was annotated and at least 95% of that annotated area belonged to one tissue class. The resulting task classifies high-purity annotated regions and does not constitute dense pixel-level segmentation of the complete WSI.

C. Data Populations and Partitions

The 513 eligible WSIs were partitioned at WSI level into three disjoint subsets: 361 for unsupervised VAE development, 129 for supervised classifier development, and 23 for a shared sealed test (Figure 4). VAE development used 322 training and 39 validation WSIs, contributing 300,000 and 30,000 unlabeled patches. Patch sampling supplied 60,000 training and 6,000 validation patches from each WSI diagnosis; diagnosis was not a VAE training target. WSI counts were not balanced by diagnosis (Table I). Supervised development used 106 training and 23 validation WSIs. Together with the shared test, these formed the 152-WSI cohort with supplemental masks. These assignments were fixed before VAE or classifier outputs on that cohort were inspected. They were shared by

both downstream tasks and used diagnosis and annotated-tissue availability, but not image pixels, VAE representations or model predictions. The same 23 test WSIs were also used for reconstruction evaluation.

For WSI diagnosis, each slide was one labeled MIL bag containing all its foreground patches selected by Otsu thresholding. Each complete WSI contributed one diagnosis label shared by all patches in its bag. The patch-level tissue classifier instead used high-purity annotated tumor, stroma and necrosis patches. Its training subsets were balanced at the five nested label budgets introduced above, whereas validation and test used fixed, unbalanced patch sets. Table I separates the inputs and class counts of these three stages.

D. Experimental Design

The same frozen VAE pair was evaluated in three tasks. First, reconstruction measured whether each encoder–decoder recovered clean, unseen patches. Second, a patch classifier used frozen latent maps to predict local tissue type. It was trained at five nested label budgets, with the same validation and test sets at every budget. Third, a multiple-instance learning (MIL) classifier aggregated patch embeddings to predict one diagnosis per WSI.

The VAEs were trained on unlabeled patches and frozen before classification. Reconstruction used each encoder and decoder; classification used only the deterministic posterior mean. Both VAEs received the same ordered patches, normalization, corruption policy, loss, update count and validation schedule. Downstream pairs shared WSI assignments, patch coordinates, classifier initialization, epoch order and model-selection rule. The architectures differ in parameter count and computational cost (Table II), which also matter when interpreting outcomes.

E. Variational Autoencoders and Shared Architecture

An autoencoder learns to turn an image into a smaller representation and then reconstruct the image from it. Its *encoder* extracts and compresses visual features, and its *decoder* uses the compressed embedding to reproduce the input. In a VAE,




Fig. 3. Images and annotations in UBC-OCEAN [10], [19]. Panels (a)–(c) show a WSI thumbnail and two actual 256×256 crops from that WSI; arrows mark their source regions. Panels (d)–(f) show a different, paired WSI thumbnail, supplemental tissue mask and overlay. Red marks tumor, green stroma and blue necrosis when present; black mask pixels are unannotated. Panel (g) shows a TMA. The WSI in (d) measures 40063×42207 pixels, compared with 2964×2964 for the TMA. The panels do not share a physical scale; only WSIs enter the experiment.

the encoder specifies a distribution over embeddings rather than one fixed code: here a diagonal Gaussian with mean μ and log-variance $\log \sigma^2$. During training, a latent sample $\mathbf{z}$ is drawn by reparameterization [11]. The distribution is regularized toward a unit Gaussian while the decoder learns image reconstruction. After training, the posterior mean is a deterministic feature map for the classifiers. For reconstruction evaluation, the decoder reconstructs the clean patch from that mean. Figure 5 shows these paths.

Both models used a four-stage residual encoder informed by ResNet-18 [21]: a stem and two residual blocks at each stage, with $2\times$ fixed anti-aliased downsampling between stages. Their decoder used two residual blocks per stage and reversed the three resolution changes with bilinear upsampling followed by convolution. The spatial sequence was $256 \rightarrow 128 \rightarrow 64 \rightarrow 32$ pixels per side in the encoder and reversed in the decoder. Separate posterior heads produced

$$\mu, \log \sigma^2 \in \mathbb{R}^{16 \times 32 \times 32}. \quad (1)$$

The decoder returned a raw three-channel image from a zero-initialized output convolution, without a final tanh. The 16-channel spatial embedding preserves coarse locations within

the patch; its 32×32 grid is smaller than the original 256×256 input. Table II compares the two implementations.

F. Continuous Rotation Equivariance

A *group* is a collection of transformations that can be composed and inverted. The planar rotation group $\text{SO}(2)$ contains every angle from 0 to 360 degrees, with angle addition as its composition rule. A map f is *equivariant* when transforming its input also transforms its output in a predictable way:

$$f(T_g \mathbf{x}) = \rho(g)f(\mathbf{x}), \quad g \in \text{SO}(2). \quad (2)$$

Here T_g rotates an image and $\rho(g)$ specifies how the output feature field must rotate. This is distinct from *invariance*, in which the output stays identical after rotation. For a scalar image feature, spatial locations rotate but the value at a corresponding location does not; for a vector feature, its two components rotate as well.

The conventional VAE used ordinary convolutions, Group-Norm and learned scalar gates. Its hidden channel counts were 32, 48, 64 and 96. The experimental $\text{SO}(2)$ VAE replaced the convolutional layers with steerable convolutions [5], [6]. A steerable kernel is a filter whose rotated version follows




Fig. 4. Data populations and WSI-level partitions. The source cohort is split into unsupervised VAE development, supervised classifier development, and a separate shared sealed test. The same 23 test WSIs support all three final evaluations, with task-specific patch populations. All assignments are fixed at WSI level and shared by both representations. Numbers inside split boxes follow the order CC/EC/HGSC/LGSC/MC.

TABLE I
INPUTS BY TASK AND SPLIT. DIAGNOSIS ORDER IS CC/EC/HGSC/LGSC/MC; TISSUE ORDER IS TUMOR/STROMA/NECROSIS. VAE PATCH DIAGNOSIS IS INHERITED FROM THE SOURCE WSI, NOT A LOCAL TISSUE LABEL. MIL CLASS COUNTS REFER TO WSIs; TISSUE CLASS COUNTS REFER TO PATCHES.

Task	Split	WSIs	Patches	Class composition
VAE	Train	322	300,000	54/74/144/24/26 WSIs; 60,000 patches per diagnosis.
	Validation	39	30,000	7/9/17/3/3 WSIs; 6,000 patches per diagnosis.
Reconstruction	Test	23	67,138	Foreground patches; uncertainty aggregated by WSI.
Tissue	Train	106	750–17,013	250–5,671 selected patches per tissue class.
	Validation	23	31,339	22,255/7,809/1,275 patches.
	Test	23	31,572	21,796/8,500/1,276 patches.
WSI MIL	Train	106	1,224,875	23/24/40/11/8 WSI labels; complete foreground bags.
	Validation	23	264,178	5/6/8/2/2 WSI labels; complete foreground bags.
	Test	23	261,168	5/6/8/2/2 WSI labels; complete foreground bags.

a prescribed rule: rather than learning an unrelated filter for every orientation, the network learns coefficients of radial and angular basis functions that obey the input/output field transformation laws. This construction constrains the convolutional operators; it does not by itself establish that the complete trained encoder obeys Equation 2 exactly under sampled image rotations.

Its hidden features contained F_0 scalar fields and F_1 vector fields. An F_0 copy stores one value per pixel and changes only its spatial location under a rotation. An F_1 copy stores an ordered pair of values per pixel; that pair rotates by the same angle as the image coordinate system. Thus, a layout with n_0 scalar copies and n_1 vector copies has $n_0 + 2n_1$ numerical channels, not $n_0 + n_1$. The four stages used

$(n_0, n_1) = (16, 16), (24, 24), (32, 32), (48, 48)$. Field-aware normalization and gates based on the F_1 vector magnitude were used so these operations respect the field types. The precise nonlinearities and normalizations are given below. The posterior heads returned 16 F_0 fields and no F_1 fields, matching the conventional $16 \times 32 \times 32$ posterior.

1) *Kernel Support and Learned Parameterization*: Comparing conventional and rotation-equivariant CNNs is not straightforward: matching kernel support, learned parameter count and computational cost can require different architectural choices. The conventional VAE uses 5×5 residual kernels, whereas the $SO(2)$ VAE uses 7×7 steerable kernels. This difference follows from the spatial sampling requirements of the selected steerable basis.

The basis combines angular harmonics with Gaussian radial




Fig. 5. Shared denoising VAE workflow. Both models have the same input, posterior shape and encoder–decoder resolution schedule. They differ in the type of convolution and the accompanying field normalization and gates. Training reconstructs the clean patch from a possibly corrupted input; classification uses only the frozen posterior mean. The dashed rectangular path supplies the clean input as the reconstruction target.

profiles, a common construction in steerable CNNs [6]. The radial envelopes are centered at approximately 1.00, 1.90 and 2.75 pixels, each with a standard deviation of 0.30 pixels. Its F_1 -to- F_1 terms include a second angular harmonic, corresponding to two cycles around the filter center. A centered 5×5 grid extends only two pixels along each coordinate axis; only its four corners lie beyond radius 2.5 pixels. Consequently, the outer radial envelope is sampled unevenly across angles. The 7×7 grid extends three pixels along each axis and provides additional samples near this envelope, improving its angular coverage and reducing the risk of discretization aliasing.

While the larger grid might appear to give the steerable layers an advantage, their filters remain constrained to a fixed steerable basis. They learn coefficients of this basis rather than an independent weight at every kernel location. For example, the steerable convolution in the first same-resolution residual stage uses only 30.0% as many learned parameters as its conventional counterpart: 7,680 basis coefficients versus 25,600 free weights. Across the complete VAEs, the SO(2) model has 29.8% as many learned parameters, although its estimated convolutional work is higher (Table II).

Larger steerable than conventional kernels have also been used in earlier comparisons. Weiler and Cesa used larger steerable kernels to mitigate discretization artifacts while adjusting feature widths to preserve the total parameter count [6]. Worrall *et al.* compared 5×5 harmonic filters with 3×3 conventional filters and reported their different parameter counts [22]. Figure 6 shows the two residual building blocks before Figure 7 assembles them into the full encoder and decoder schedules.

G. Field Gates and Normalization

The conventional VAE applies a learned scalar gate to each hidden channel, including after residual addition. This gate is a parameterized variant of the sigmoid-weighted linear unit (SiLU), $u \sigma(u)$ [23], with a learned slope and offset inside the sigmoid. For a channel value u_c ,

$$A_0(u_c) = u_c \sigma(a_c u_c + b_c), \quad \sigma(t) = \frac{1}{1 + e^{-t}}, \quad (3)$$

where a_c and b_c are learned per-channel parameters, initialized to $a_c = 1$ and $b_c = 0$, so the initial activation is SiLU. The F_0 fields in the SO(2) VAE use the same rule with one parameter pair per scalar field. For an F_1 field $\mathbf{v}_j = (v_{j,1}, v_{j,2})$, the

equivariant VAE gates both components by the same factor computed from their rotation-invariant magnitude:

$$r_j = \sqrt{v_{j,1}^2 + v_{j,2}^2 + 10^{-4}}, \quad (4)$$

$$A_1(\mathbf{v}_j) = \mathbf{v}_j \sigma(\alpha_j r_j + \beta_j).$$

The learned α_j, β_j are shared by the two components. A planar rotation changes the vector direction but not r_j , so this gate rotates its output together with the input vector.

Before these gates, the conventional network uses affine GroupNorm with eight groups. For each sample and channel c , it computes the mean and variance over the channels in c 's group and all spatial positions:

$$N_{\text{GN}}(u_c) = \gamma_c \frac{u_c - \mu_{G(c)}}{\sqrt{s_{G(c)}^2 + 10^{-5}}} + \delta_c. \quad (5)$$

The equivariant network applies this same eight-group formula only to its F_0 scalar fields. Its F_1 vectors instead use four groups and a shared root-mean-square denominator for both components:

$$N_1(\mathbf{v}_j) = \gamma_j \frac{\mathbf{v}_j}{\sqrt{\text{mean}_{k \in G(j), d \in \{1,2\}, p} v_{k,d,p}^2 + 10^{-5}}}. \quad (6)$$

Here p indexes spatial positions within one sample. The vector rule has no mean subtraction or additive bias, because either would introduce a preferred direction. Its denominator is unchanged by rotations in feature space. The normalization statistics and gates are evaluated in FP32 in the SO(2) implementation.

H. VAE Objective and Training

Both models used a denoising VAE training criterion: an encoder sees a possibly corrupted patch, and the decoder is trained to reconstruct the corresponding clean patch [12]. In the implemented training procedure, with probability 0.3 the corruption operator C perturbed hematoxylin/eosin/residual stain coordinates and added image-space Gaussian noise [24]. The resulting input was clamped to the normalized $[-1, 1]$ RGB domain. A patch could therefore be clean or corrupted at the input, while its reconstruction target was always the same clean patch:

$$\mathbf{x} \xrightarrow{C} \tilde{\mathbf{x}} \xrightarrow{q_\phi} \mathbf{z} \xrightarrow{D_\theta} \hat{\mathbf{x}} \text{ compared with } \mathbf{x}. \quad (7)$$





Fig. 6. Same-resolution residual blocks used by the conventional and continuous-SO(2) VAEs. Each main branch contains two learned convolutions; normalization and a gate follow the first convolution, while the second is followed by normalization. The identity skip is added before the output gate. At encoder stage transitions, both branches are downsampled separately with a fixed 5×5 binomial filter: on the main branch after the first gate and on the skip branch before its learned projection and normalization. At decoder transitions, both branches are bilinearly upsampled before their first learned convolution; the skip branch again uses a projection and normalization. The F_0/F_1 field types and gate equations are defined in the surrounding text. The 7×7 steerable support provides fuller sampling of the selected radial and angular basis; Section III-F1 gives the rationale and learned-parameter comparison.

Encoder paths



Decoder paths




Fig. 7. Layer-level architecture of the two frozen VAEs. The residual building blocks are detailed in Figure 6. Each stage A–D contains two residual blocks; stage-changing blocks additionally resample both branches and project the skip branch. A conventional hidden width is a channel count; in the SO(2) model, $nF_0 + nF_1$ means n scalar fields and n two-component vector fields. Posterior mean and log-variance use separate heads; training samples $\mathbf{z}$, whereas downstream classifiers use μ . The conventional stem uses a 7×7 kernel and its other convolutions use 5×5 kernels; the SO(2) stem uses a 9×9 steerable kernel and its other convolutions use 7×7 kernels. See Section III-F1 for the kernel-support rationale and learned-parameter comparison.

The encoder distribution $q_\phi(\mathbf{z} | \tilde{\mathbf{x}})$ is the diagonal Gaussian determined by the mean and log-variance heads; D_θ is the decoder. During training, $\mathbf{z} = \mu + \sigma \odot \epsilon$, with $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$. The implemented loss was

$$\mathcal{L} = \text{MAE}(\hat{\mathbf{x}}, \mathbf{x}) + 0.1[1 - \text{SSIM}(\hat{\mathbf{x}}, \mathbf{x})] + \beta D_{\text{KL}}(q_\phi(\mathbf{z} | \tilde{\mathbf{x}}) \| \mathcal{N}(\mathbf{0}, \mathbf{I})), \quad (8)$$

where MAE measures pixelwise reconstruction error, SSIM measures structural similarity [25], and the KL term regularizes the input-dependent posterior toward a unit Gaussian. The KL weight increased linearly from zero to 0.01 during the first 6,000 updates and remained fixed thereafter.

AdamW [26] used coefficients (0.9, 0.999), $\epsilon = 10^{-8}$, weight decay 10^{-5} and gradient-norm clipping at 1.0. The learning rate increased linearly for 600 updates to 10^{-3} and then decreased with a cosine schedule to 10^{-5} . Each model was trained for 60,000 successful updates with a global batch of 50 patches distributed over two NVIDIA T4 GPUs. Validation and checkpointing occurred every 3,000 updates. The comparison used the final update-60,000 checkpoint for both models rather than selecting different VAE checkpoints from the downstream results.

After VAE training, all encoder and decoder parameters were frozen. Downstream tasks used the posterior mean μ

TABLE II
ARCHITECTURE COMPARISON.

Property	Conventional VAE	Continuous-SO(2) VAE
Hidden representation	Ordinary channels: 32, 48, 64, 96	(F_0, F_1) copies: (16,16), (24,24), (32,32), (48,48); 48, 72, 96, 144 numerical channels
Learned convolutions	43 ordinary convolutions	43 steerable convolutions
Normalization/nonlinearity	Eight-group GroupNorm; learned scalar gate	Eight-group F_0 GroupNorm; four-group F_1 RMS normalization; scalar and radial gates
Posterior	$\mu, \log \sigma^2 \in \mathbb{R}^{16 \times 32 \times 32}$	Same spatial Gaussian posterior, represented by scalar fields
Learned parameters	3,958,435	1,180,035
Estimated convolutional work	3.65×10^{10} MAC/sample	1.598×10^{11} MAC/sample
Rotation constraint	None beyond translation-equivariant convolution	Continuous SO(2) kernel constraint

rather than a sampled latent. For every selected coordinate, both encoders received the same RGB patch. The resulting paired latent datasets had identical WSI membership, tissue labels and patch coordinates.

I. Reconstruction Evaluation

The reconstruction test contained 67,138 foreground patches from the 23 test WSIs. For each slide, eligible foreground coordinates were sorted by row and column (y , then x). If there were more than 3,000, 3,000 evenly spaced rank positions were selected across that ordered list; otherwise all were retained. Thus, 22 slides contributed 3,000 patches each and one contributed its available 1,138. The cap limited the influence of slides with many more eligible patches. MAE and mean squared error (MSE) were measured in the normalized $[-1, 1]$ domain. Peak signal-to-noise ratio (PSNR) and SSIM were measured after conversion to $[0, 1]$ and clamping. Mean MAE was the primary reconstruction measure. The other metrics describe complementary error properties.

J. Classification A: Tissue Recognition and Label Budgets

Tissue classification uses one patch at a time, without combining information across a WSI. It tests how much tissue information each frozen representation retains as the number of labeled training patches changes. The task predicted tumor, stroma or necrosis from a single latent map. The classifier contained three convolutional stages $16 \rightarrow 32 \rightarrow 64 \rightarrow 128$, GroupNorm, GELU activations [27], adaptive average pooling and a linear $128 \rightarrow 3$ output layer.

Training subsets contained 250, 500, 1,000, 2,500 and 5,671 labeled patches per class. They were class balanced, nested and selected without replacement; every smaller set was a subset of each larger one. WSI-stratified ordering ensured that the smallest subset did not collapse onto a single slide. Each normal-equivariant pair started from identical classifier weights and received the same patch order. Models were optimized for at most 30 epochs, and the best validation macro-F1 checkpoint was retained. The test distribution was not balanced and contained 21,796 tumor, 8,500 stroma and 1,276 necrosis patches from 23 WSIs.

K. Classification B: Five-Class WSI Diagnosis

The diagnostic task predicted the five ovarian-cancer subtypes. Every selected foreground patch was represented by its frozen posterior mean. A trainable convolutional patch encoder compressed the $16 \times 32 \times 32$ latent map. Two local attention blocks exchanged information within a radius-two spatial neighborhood, and a global stream summarized the complete WSI through learned sigmoid attention. A final classification layer produced five logits. The supervised model contained approximately 1.51 million parameters.

Training used one complete WSI bag per optimizer step, AdamW and class-weighted cross-entropy computed from the 106 training WSIs. Each representation branch started from the same classifier weights and received the same WSI order. The maximum was 150 epochs, validation was performed twice per epoch, and each branch independently retained its highest validation macro-F1 checkpoint. The primary test metric was macro-F1; balanced accuracy, accuracy, mean cross-entropy, classwise F1 and confusion matrices were secondary summaries.

L. Statistical Analysis

All difference intervals are reported as conventional-minus-equivariant. Reconstruction differences used 10,000 paired cluster-bootstrap replicates at WSI level [28]. WSI diagnosis used 10,000 paired bootstrap replicates stratified by diagnosis. Tissue recognition used a paired WSI bootstrap stratified by observed tissue support. The five label-budget macro-F1 differences and the area under the log-budget learning curve (AULC) difference were considered together as six related comparisons. Their 95% intervals were adjusted simultaneously: across the six comparisons, the bootstrap procedure targets 95% coverage for the whole set of differences. A difference interval entirely below zero favors the equivariant model when higher is better; an interval entirely above zero favors the conventional model. An interval containing zero does not establish a statistically significant difference at the 5% level, nor does it establish equivalence. Patch standard deviations describe variation among patches, while confidence intervals resample WSIs as the biological units.

TABLE III
RECONSTRUCTION ON 67,138 TEST PATCHES. VALUES ARE MEAN
(POPULATION STANDARD DEVIATION).

Metric	Conventional	SO(2)
MAE ↓	0.0629 (0.0239)	0.0642 (0.0223)
MSE ↓	0.0080 (0.0074)	0.0083 (0.0067)
PSNR dB ↑	27.923 (2.721)	27.658 (2.595)
SSIM ↑	0.7575 (0.0842)	0.7436 (0.0862)

IV. RESULTS

A. VAE Reconstruction

Figure 8 displays the first nine examples from the fixed 25-patch validation selector, in selector order. Each original is paired with both frozen models’ final reconstruction.

The conventional VAE had lower pooled MAE and MSE and higher pooled PSNR and SSIM (Table III). Mean MAE was 0.062872 for the conventional model and 0.064230 for the equivariant model. The paired difference was -0.001358 , with a 95% interval of $[-0.002213, 0.000046]$. The interval included zero, so the primary analysis did not establish a reconstruction difference. The remaining metrics are descriptive secondary summaries.

Figure 9 shows descriptive variation among the 23 test WSIs. Each plotted value is the mean of that WSI’s selected test patches; the paired WSI-bootstrap interval above remains the inferential comparison.

B. Classification A: Three-Class Tissue Recognition

The tissue results changed with the amount of labeled training data (Figure 10 and Table IV). The equivariant representation had higher macro-F1 at 250, 500, 1,000 and 2,500 labels per class. At 500 labels per class, the macro-F1 values were 0.575 and 0.511, and the simultaneous interval for conventional-minus-equivariant was $[-0.121, -0.008]$. This was the only label budget whose simultaneous interval excluded zero.

At 500 labels per class, the classwise F1 values for conventional/equivariant representations were 0.769/0.781 for tumor, 0.485/0.560 for stroma and 0.279/0.384 for necrosis. These values show how the aggregate macro-F1 difference was distributed across the three tissue classes.

At the largest budget, the conventional representation obtained macro-F1 0.766 compared with 0.710 for the equivariant representation. The simultaneous interval included zero, with a lower bound of -0.001 . The corresponding classwise F1 values were 0.868/0.871 for tumor, 0.728/0.711 for stroma and 0.701/0.548 for necrosis. The normalized area under the log-budget learning curve was 0.615 for the conventional representation and 0.631 for the equivariant representation; its simultaneous difference interval included zero.

C. Classification B: Five-Class WSI Diagnosis

Macro-F1 was 0.483 for the equivariant representation and 0.389 for the conventional representation (Table V). The conventional-minus-equivariant difference was -0.095 , with

TABLE IV
TISSUE MACRO-F1 ON 31,572 TEST PATCHES. SIMULTANEOUS INTERVALS
ARE REPORTED FOR CONVENTIONAL-MINUS-SO(2).

Labels/class	Conventional ↑	SO(2) ↑	Simultaneous 95% CI
250	0.508	0.523	$[-0.072, 0.041]$
500	0.511	0.575	$[-0.121, -0.008]$
1,000	0.582	0.598	$[-0.072, 0.041]$
2,500	0.701	0.714	$[-0.069, 0.044]$
5,671	0.766	0.710	$[-0.001, 0.112]$

TABLE V
OVERALL WSI DIAGNOSIS RESULTS ON 23 TEST SLIDES. DIFFERENCE
INTERVALS ARE CONVENTIONAL-MINUS-SO(2).

Metric	Conventional	SO(2)	Difference 95% CI
Macro-F1 ↑	0.389	0.483	$[-0.207, 0.050]$
Balanced accuracy ↑	0.510	0.503	$[-0.167, 0.180]$
Accuracy ↑	0.522	0.609	$[-0.217, 0.043]$
Mean cross-entropy ↓	1.323	0.993	$[0.135, 0.524]$

a 95% interval of $[-0.207, 0.050]$. The interval included zero. Among the secondary metrics, the equivariant representation had lower mean cross-entropy, with a WSI-resampling interval excluding zero. The macro-F1 interval still crossed zero.

Classwise point estimates varied (Table VI). The conventional representation had zero F1 for EC and LGSC and correctly classified both MC slides; the equivariant representation had higher EC and LGSC F1 and no correct MC prediction. With only two test WSIs in each of LGSC and MC, these contrasts are unstable.

V. DISCUSSION

Which representation scored higher depended on the task and the number of labeled tissue patches. The conventional VAE had lower pooled MAE/MSE and higher pooled PSNR/SSIM point estimates. The equivariant representation had higher tissue macro-F1 at the four smaller budgets and a higher WSI macro-F1 point estimate. Most confidence intervals for the main comparisons included zero. The single tissue-budget interval that excluded zero applies to that contrast only and does not establish a monotone or general label-efficiency effect. The reconstruction point-estimate ordering was not consistent with the downstream point-estimate ordering.

Patch-based WSI systems depend on both the instance representation and the bag aggregator. Prior MIL work has developed the representation, aggregation or both [15]–[18]. This experiment held the downstream protocol fixed while changing the frozen representation. The WSI interval remained wide, and classwise estimates were unstable for diagnoses with two test slides. The results characterize the two implemented representation branches but do not establish a general WSI or tissue-classification advantage.

The models also differ in capacity and computation. Graham *et al.* matched parameter counts in comparisons of rotation-equivariant histology classifiers [14]; these two VAEs were not parameter-matched. The equivariant VAE uses 29.8%




Fig. 8. Three matched 3×3 mosaics: original patches, conventional-VAE reconstructions and continuous-SO(2)-VAE reconstructions. The nine examples are the first nine of the fixed validation selector; quantitative results use 67,138 patches from 23 test WSIs.




Fig. 9. Reconstruction metric distributions across 23 test WSIs. Each point is one WSI-level mean; boxes summarize those 23 means for each frozen VAE. They are descriptive distributions, not confidence intervals. MAE and MSE use normalized RGB; PSNR and SSIM use image-domain values.

TABLE VI
CLASSWISE WSI F1 AND TEST SUPPORT.

Diagnosis	Support	Conventional F1 $\uparrow$	SO(2) F1 $\uparrow$
CC	5	0.667	0.750
EC	6	0.000	0.533
HGSC	8	0.706	0.800
LGSC	2	0.000	0.333
MC	2	0.571	0.000

as many learned parameters but requires more estimated multiply-accumulate operations. Any performance difference

may reflect equivariance, parameter count, optimization or their interaction. The experiment does not separate these mechanisms.

VI. LIMITATIONS

Each VAE and downstream configuration was trained once. The bootstrap intervals quantify variation across the observed WSIs, not variation across new initializations, data orders or optimization runs. The WSI test contains 23 slides, including two each for LGSC and MC. Tissue recognition contains many patches but only 23 WSIs, and necrosis is present in five of them.




Fig. 10. Frozen-representation classification results. (a) Tissue macro-F1 on 31,572 test patches under nested class-balanced training budgets; the asterisk marks the 500-label contrast whose simultaneous 95% difference interval excludes zero. (b) Five-class diagnosis on 23 test WSIs. The paired WSI macro-F1 interval includes zero.

Supplemental tissue masks are incomplete. The tissue experiment evaluates selected high-purity annotated regions and cannot be interpreted as exhaustive segmentation of complete WSIs. The study uses one ovarian-cancer data source.

VII. CONCLUSION

This study compared conventional and continuous-SO(2) VAEs as frozen histopathology representations under a common experimental protocol. With one training run per configuration and intervals limited to WSI resampling, the results varied by task and label budget. The experiment did not establish a general advantage for either implemented representation; repeated training runs and models with matched parameter counts are needed to isolate the contribution of continuous rotation equivariance.

REFERENCES

- [1] G. Litjens, T. Kooi, B. E. Bejnordi, A. A. A. Setio, F. Ciompi, M. Ghafoorian, J. A. W. M. van der Laak, B. van Ginneken, and C. I. Sanchez, "A survey on deep learning in medical image analysis," *Medical Image Analysis*, vol. 42, pp. 60–88, 2017.
- [2] A. H. Song, G. Jaume, D. F. K. Williamson, M. Y. Lu, A. Vaidya, T. R. Miller, and F. Mahmood, "Artificial intelligence for digital and computational pathology," *Nature Reviews Bioengineering*, vol. 1, no. 12, pp. 930–949, 2023.
- [3] D. Tellez, G. Litjens, J. van der Laak, and F. Ciompi, "Neural image compression for gigapixel histopathology image analysis," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 43, no. 2, pp. 567–578, 2021.
- [4] T. S. Cohen and M. Welling, "Group equivariant convolutional networks," in *Proceedings of the 33rd International Conference on Machine Learning*, 2016, pp. 2990–2999.
- [5] T. S. Cohen and M. Welling, "Steerable CNNs," in *International Conference on Learning Representations*, 2017.
- [6] M. Weiler and G. Cesa, "General $E(2)$ -equivariant steerable CNNs," in *Advances in Neural Information Processing Systems*, vol. 32, 2019.
- [7] B. S. Veeling, J. Linmans, J. Winkens, T. Cohen, and M. Welling, "Rotation equivariant CNNs for digital pathology," in *Medical Image Computing and Computer Assisted Intervention – MICCAI 2018*. Springer, 2018, pp. 210–218.
- [8] A. Nasiri and T. Bepler, "Unsupervised object representation learning using translation and rotation group equivariant VAE," in *Advances in Neural Information Processing Systems*, vol. 35, 2022.
- [9] J. Burgess, J. J. Nirschl, M.-C. Zanellati, A. Lozano, S. Cohen, and S. Yeung-Levy, "Orientation-invariant autoencoders learn robust representations for shape profiling of cells and organelles," *Nature Communications*, vol. 15, p. 1022, 2024.
- [10] University of British Columbia, "UBC ovarian cancer subtype classification and outlier detection (UBC-OCEAN)," Kaggle competition, 2023. [Online]. Available: <https://www.kaggle.com/competitions/UBC-OCEAN>
- [11] D. P. Kingma and M. Welling, "Auto-encoding variational bayes," in *International Conference on Learning Representations*, 2014.
- [12] D. J. Im, S. Ahn, R. Memisevic, and Y. Bengio, "Denoising criterion for variational auto-encoding framework," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 31, no. 1, 2017, pp. 2059–2065.

- [13] B. Elesedy and S. Zaidi, "Provably strict generalisation benefit for equivariant models," in *Proceedings of the 38th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 139, 2021, pp. 2959–2969.
- [14] S. Graham, D. Epstein, and N. Rajpoot, "Dense steerable filter CNNs for exploiting rotational symmetry in histology images," *IEEE Transactions on Medical Imaging*, vol. 39, no. 12, pp. 4124–4136, 2020.
- [15] M. Ilse, J. M. Tomczak, and M. Welling, "Attention-based deep multiple instance learning," in *Proceedings of the 35th International Conference on Machine Learning*, vol. 80, 2018, pp. 2127–2136.
- [16] M. Y. Lu, D. F. K. Williamson, T. Y. Chen, R. J. Chen, M. Barbieri, and F. Mahmood, "Data-efficient and weakly supervised computational pathology on whole-slide images," *Nature Biomedical Engineering*, vol. 5, pp. 555–570, 2021.
- [17] B. Li, Y. Li, and K. W. Eliceiri, "Dual-stream multiple instance learning network for whole slide image classification with self-supervised contrastive learning," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2021, pp. 14 318–14 328.
- [18] Z. Shao, H. Bian, Y. Chen, Y. Wang, J. Zhang, X. Ji, and Y. Zhang, "TransMIL: Transformer based correlated multiple instance learning for whole slide image classification," in *Advances in Neural Information Processing Systems*, vol. 34, 2021, pp. 2136–2147.
- [19] sohier, "UBC Ovarian Cancer Competition Supplemental Masks," Kaggle dataset, 2023. [Online]. Available: <https://www.kaggle.com/datasets/sohier/ubc-ovarian-cancer-competition-supplemental-masks>
- [20] N. Otsu, "A threshold selection method from gray-level histograms," *IEEE Transactions on Systems, Man, and Cybernetics*, vol. 9, no. 1, pp. 62–66, 1979.
- [21] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2016, pp. 770–778.
- [22] D. E. Worrall, S. J. Garbin, D. Turmukhambetov, and G. J. Brostow, "Harmonic networks: Deep translation and rotation equivariance," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2017.
- [23] S. Elfving, E. Uchibe, and K. Doya, "Sigmoid-weighted linear units for neural network function approximation in reinforcement learning," *Neural Networks*, vol. 107, pp. 3–11, 2018.
- [24] A. C. Ruifrok and D. A. Johnston, "Quantification of histochemical staining by color deconvolution," *Analytical and Quantitative Cytology and Histology*, vol. 23, no. 4, pp. 291–299, 2001.
- [25] Z. Wang, A. C. Bovik, H. R. Sheikh, and E. P. Simoncelli, "Image quality assessment: From error visibility to structural similarity," *IEEE Transactions on Image Processing*, vol. 13, no. 4, pp. 600–612, 2004.
- [26] I. Loshchilov and F. Hutter, "Decoupled weight decay regularization," in *International Conference on Learning Representations*, 2019.
- [27] D. Hendrycks and K. Gimpel, "Gaussian error linear units (GELUs)," 2016.
- [28] C. A. Field and A. H. Welsh, "Bootstrapping clustered data," *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, vol. 69, no. 3, pp. 369–390, 2007.